


CodexClaw Full Gap Strategy Report

A governance-first thin harness can close product, operating, and engineering gaps through five focused programs.

Reference	Pinned scope
CodexClaw	98cac96 · local current source
OMO beta	84f98d8 · v5.0.0-beta.22
Senpi	703d9d7 · 2026.8.27 (OMO peer 2026.8.26-2)
Coverage	5 pillars · 20 domains · 42 pages

CodexClaw should scale governance, not recreate the runtime

The product's advantage is durable policy and evidence over native Codex; the next stage is integration, operator experience, and measured delivery.



Five programs close the largest gaps without breaking the thin-host boundary

The executive thumbnail states what to fund; page 34 later tests the portfolio against impact and feasibility.






Exhibit 2 · Evidence readout

Sources C2, C10, C12, C16, C17. Decision sequence – P1: Control-plane truth → P2: Operator experience → P3: Durable operations → P4: Static safety → P5: Distribution & security

Management implication

Fund five programs in dependency order rather than pursuing competitor feature parity.



Evidence supports directional maturity, not benchmark rankings

This executive caveat states how to read the scores; Appendix A later provides the auditable method and source ledger.

Evidence base

C5, C16, O7, S9

Diagnostic question

Bound how scores may be used.

Decision test

Merge to appendix method only if removed from executive section.

Pinned source

Runtime vs prompt

Integer stage

Confidence

Decision

Exhibit 3 · Evidence readout

Sources C5, C16, O7, S9. Decision sequence – Pinned source → Runtime vs prompt → Integer stage → Confidence → Decision

Management implication

Use the heatmap to prioritize domains, not to publish a product winner.

Accountable owner

transformation lead + program owners

Primary KPI

five program exit gates, reviewed monthly

Principal risk

portfolio dilution or runtime duplication

Five MECE pillars define what 'full gap' means

Twenty domains cover strategy, assurance, product operations, engineering platform, and scale/ecosystem without double-counting owners.






Exhibit 4 · Evidence readout

Sources C1, C5, C10, C16, C17. Decision sequence – A: Strategy & execution → B: Assurance & trust → C: Product & operations → D: Engineering platform → E: Scale & ecosystem

Management implication

A page only survives if it changes priority, owner, KPI, transfer, or gate.



CodexClaw owns policy and evidence; Codex owns execution

The owner boundary is the central strategic differentiator and the filter for every transfer decision.





Exhibit 5 · Evidence readout

Sources C1, C4, O4, S7. Decision sequence – Prompt: methodology → Plugin: policy & evidence → Host: execution runtime

Management implication

Do not create a second scheduler, terminal registry, goal DB, or Stop owner.



Maturity is bimodal: governance is strong, product operations are fragmented

CodexClaw scores 4 in orchestration/evidence/release, but 2 in console, observability, routing, workspace intelligence, maintainability, and performance operations.

Evidence base
C1, C5, C10, C12, C16, C19

Diagnostic question
Identify portfolio imbalance.

Decision test
Merge page 8 if one heatmap survives.

	Orchestration	Research	Evidence	UX	Observability	Release	Maintainability	Ecosystem
CodexClaw	4	3	4	3	2	4	2	3
OMO	4	4	4	4	4	4	3	4
Senpi	3	2	3	N/E	4	4	4	3

Exhibit 6 · Evidence readout
Sources C1, C5, C10, C12, C16, C19. CodexClaw: 4 in orchestration/evidence/release; 2 in observability/maintainability. Senpi UX is N/E, not zero.

Management implication
Balance the control plane with product operations and engineering discipline.

Accountable owner
transformation lead + program owners

Primary KPI
five program exit gates, reviewed monthly

Principal risk
portfolio dilution or runtime duplication

Sources: C1, C5, C10, C12, C16, C19 ·

CodexClaw wins where truth is durable and loses where state is disconnected

The causal gap is not missing features; it is weak connection between existing PABCD, bridge, capability, evidence, and delivery state.






Exhibit 7 · Evidence readout

Sources C2, C5, C10, C12. Decision sequence – Durable truth: FSM/goalplan/receipts → Disconnected views: GUI/bridge/capabilities → Management action: connect, measure, gate

Management implication

Connect authoritative state before adding new modes or agents.



PABCD is real runtime, but intellectual quality remains partly prompt-bound

Legal transitions and receipts are mechanized; dependency slicing, verifier realism, and adversarial reasoning still depend on model compliance.

Evidence base

C1, O1

Diagnostic question

Tier governance by work risk.

Decision test

Merge page 11 only if FSM evidence stays visible.

A. Strategy & execution maturity (1–5; N/E excluded)

CodexClaw

4

OMO

4

Senpi

3

Exhibit 8 · Evidence readout

Sources C1, O1. Directional maturity – CodexClaw 4; OMO 4; Senpi 3. Shipped: FSM, plan gate, check gate

Management implication

Keep lean flow for C0–C2; require bound review and stronger evidence for C3/C4.

Evidence-backed diagnosis

- Shipped: FSM, plan gate, check gate
- Prompt: plan quality, synthesis, reviewer reuse

Accountable owner

pabcd-state

Primary KPI

0 unsupported transitions

Principal risk

ceremony ≠ quality

Goalplan blocks false completion but cannot operate the full task lifecycle

Tasks and criteria affect completion, yet public add-task, resolve-task, and evidence-capture operations do not exist.





Exhibit 9 · Evidence readout


Sources C2, O2, S1. Directional maturity – CodexClaw 3; OMO 4; Senpi 4. Current: pending/done checklist

Management implication

Extend the existing authoritative goalplan; do not add a second task database.

Evidence-backed diagnosis

- Current: pending/done checklist
- Needed: owner, outcome, evidence, retry lineage



Sources: C2, O2, S1 ·

Review identity is strong when bound and form-only when lean

Launch IDs, plan hashes, and reviewer identity are durable only when a review round is opened; lean attestation remains unauthenticated text.





Exhibit 10 · Evidence readout

Sources C1, C5, O5. Decision sequence – C0-C2: lean attest → C3-C4: bound review round → Final gate: source + reviewer + QA

Management implication

Require current-plan review binding for C3/C4 while preserving bounded recovery.



Research quality rules outpace the state that must recover them

Tier 3 specifies waves, leads, and claims, but compaction cannot deterministically recover the next research action.

Evidence base C3, O3	Diagnostic question Fund durable research projection.	Decision test Keep if distinct research consumer case.
--------------------------------	---	--

A. Strategy & execution maturity (1–5; N/E excluded)




Exhibit 11 · Evidence readout

Sources C3, O3. Directional maturity – CodexClaw 3; OMO 4; Senpi 2. CodexClaw: proof ladder, no durable wave state

Management implication

Map run → work phase, lane/lead → task, claim → criterion, ledger → history.


Evidence-backed diagnosis

- CodexClaw: proof ladder, no durable wave state
- OMO: deep claim/observation lineage

Accountable owner search + goalplan	Primary KPI >=95% lane disposition	Principal risk unused claim ceremony
---	--	--

Search proof is strong; capability-aware routing is weak

Discovery-versus-proof doctrine is excellent, but the live search ladder is not rewritten from the tools actually available this turn.



Recall is mature retrieval, not epistemic truth maintenance

FTS, memory ranking, and compaction reinjection are shipped; durable claims still lack validity, counterevidence, and corroboration fields.

Evidence base

C3, O3

Diagnostic question

Add claim validity before memory promotion.

Decision test

Keep if distinct memory decision.

A. Strategy & execution maturity (1–5; N/E excluded)

CodexClaw

3

OMO

4

Senpi

2

Exhibit 13 · Evidence readout

Sources C3, O3. Directional maturity – CodexClaw 3; OMO 4; Senpi 2. Strength: CWD-scoped untrusted recall

Management implication

Promote only proven claims with source validity; keep retrieval and truth separate.

Evidence-backed diagnosis

- Strength: CWD-scoped untrusted recall
- Gap: stale curated memory can outrank newer evidence

Accountable owner

recall + search

Primary KPI

100% promoted claims valid

Principal risk

stale curated memory

Skill delivery is robust; external skill trust remains mutable

Mentioned built-in skills are bounded and inlined; external discovery fetches mutable bodies without immutable revision or permission receipt.

Evidence base

C4, C20, O11

Diagnostic question

Harden skill supply-chain review.

Decision test

Merge page 33 only if delivery versus trust remains distinct.

Discover

Inspect

Pin digest

Permission diff

Approve/rollback

Exhibit 14 · Evidence readout

Sources C4, C20, O11. Decision sequence – Discover → Inspect → Pin digest → Permission diff → Approve/rollback

Management implication

Pin revision/hash/license/permissions and separate review from activation.

Accountable owner

skill-search + security

Primary KPI

100% immutable approved

Principal risk

skill/MCP supply chain

Workspace intelligence is advertised but incomplete in the installed payload

SessionStart and plugin prompts advertise cxc map, while the payload dispatcher refuses the command as checkout-only.





Exhibit 15 · Evidence readout

Sources C14. Decision sequence – Installed prompt: 'run cxc map' → Payload command: refused → Decision: ship or remove claim

Management implication

Either ship a bounded payload map smoke or remove installed affordances.



Multi-agent policy is strong; runtime-surface detection can misroute

Role stores, leaf guards, and evidence gates are shipped; config-only V1/V2 inference and timeout semantics can choose the wrong lifecycle path.

Evidence base C4, C13, O4	Diagnostic question Build a session capability adapter.	Decision test Keep if top P1 initiative.
-------------------------------------	---	--

A/D. Execution & platform maturity (1–5; N/E excluded)



Exhibit 16 · Evidence readout
Sources C4, C13, O4. Directional maturity – CodexClaw 3; OMO 4; Senpi 3. Strength: bounded roles + skill transport

Management implication
Snapshot effective schema, identity, lifecycle verbs, concurrency, and unavailable capabilities before spawn.

Evidence-backed diagnosis

- Strength: bounded roles + skill transport
- Gap: surface truth and wait semantics

Accountable owner subagent-config	Primary KPI 0 misroutes / false deaths	Principal risk host schema drift
---	--	--

Sources: C4, C13, O4 ·

Host-native background execution lacks an integrated liveness projection

Codex owns workers and terminals, but CodexClaw lacks an authenticated cross-surface view of live child/terminal status.

Evidence base

C12, O4, S7

Diagnostic question

Keep scheduling host-owned and request live events.

Decision test

Merge page 27 only if scheduler rejection stays explicit.

Host child/terminal

Authenticated event API (missing)

CodexClaw Stop/console

Exhibit 17 · Evidence readout

Sources C12, O4, S7. Decision sequence – Host child/terminal → Authenticated event API: (missing) → CodexClaw Stop/console

Management implication

Consume authenticated host snapshots; never synthesize a local wake bus.

Accountable owner

host API + bridge ops

Primary KPI

100% authenticated live events

Principal risk

synthetic stale state

Source provenance leads, while receipt truth is still self-authored

Exact source identity invalidates stale evidence, but a writer can still claim a command and result that were never independently observed.






Exhibit 18 · Evidence readout

Sources C5, O5. Decision sequence – Source identity → Command/output digest → Executor identity → Independent review → Completion

Management implication

Bind command digest, output digest, exit status, executor identity, timing, and review round.



Final gates are sophisticated, but the strongest schema remains optional

Schema-v2 binds reviewer, test, QA, and source; marker deletion and fail-open paths keep it from being universal.

Evidence base

C5, O5

Diagnostic question

Mandate the strongest gate by work class.

Decision test

Merge page 19 only if optionality remains explicit.

B. Assurance & trust maturity (1–5; N/E excluded)

CodexClaw

4

OMO

4

Senpi

3

Exhibit 19 · Evidence readout

Sources C5, O5. Directional maturity – CodexClaw 4; OMO 4; Senpi 3. C0–C2: lean

Management implication

Require schema v2 for C3/C4 and release work; keep bounded escape paths visible as non-success.

Evidence-backed diagnosis

- C0–C2: lean
- C3–C4: source-bound

Accountable owner

goal/final-gate owner

Primary KPI

100% schema-v2 C3/C4

Principal risk

trap from over-fail-close

Hook trust protects local integrity, not publisher authenticity

Hash trust and atomic retrust prove installed consistency; they do not establish a signed publisher-to-payload chain.





Exhibit 20 · Evidence readout

Sources C6, S4. Decision sequence – Publisher → Signed release → Payload digest → Hook hash → Live trust


Management implication

Bind trusted hooks to SBOM, signature, release SHA, and post-install verification.



Security release assurance lags the repository's stated standard

CI/release includes tests, inventory, and packaging but lacks dedicated SAST, dependency audit, secret scan, SBOM, and signing.



B. Assurance & trust maturity (1–5; N/E excluded)



Exhibit 21 · Evidence readout


Sources C6, C7, S4. Directional maturity – CodexClaw 3; OMO 3; Senpi 4. Current: exact candidate design

Management implication

Require zero unresolved high/critical findings and signed, verified assets.

Evidence-backed diagnosis

- Current: exact candidate design
- Missing: SAST/SBOM/signing/secret/dependency proof



Sources: C6, C7, S4 ·

Remote operations is a differentiated shipped strength

Telegram/Discord pairing, commands, approval relay, progress, media, sessions, and dashboard are real runtime surfaces.

Evidence base C11	Diagnostic question Preserve and extend a unique advantage.	Decision test Keep if positive differentiator.
-----------------------------	---	--

C. Product & operations maturity (1–5; N/E excluded)

CodexClaw	4
OMO	N/E
Senpi	N/E

Exhibit 22 · Evidence readout

Sources C11. Directional maturity – CodexClaw 4; OMO N/E; Senpi N/E. Shipped: pairing, approvals, progress, sessions

Management implication

Use the bridge as the entry point for the unified operator console.

Evidence-backed diagnosis

- Shipped: pairing, approvals, progress, sessions
- N/E: no comparable bounded peer surface

Accountable owner messenger-bridge	Primary KPI >=99% terminal delivery	Principal risk remote full-access scope
--	---	---

Sources: C11 ·

The operator console omits the core PABCD product state

The dashboard is bridge-centric and does not surface phase, work phase, criteria, evidence, doctor, or hook trust.




Exhibit 23 · Evidence readout

Sources C10, C2, O8. The proposed console joins phase, open criteria, trust state, next action, and evidence navigation without creating a second state store.

Management implication

Show workflow truth and next action without creating parallel state.

Accountable owner

GUI/product owner

Primary KPI

state visibility <2s

Principal risk

parallel UI state

Sources: C10, C2, O8 ·

Installation is concise but not guided end-to-end

Two commands install the plugin, but hook trust, doctor, model choices, console, and optional bridge are not one progressive journey.






Exhibit 24 · Evidence readout

Sources C9, O8, S10. Decision sequence – Install → Approve hooks → Doctor → Choose models → First task

Management implication

Health → trust → models → optional bridge → first successful task in ≤4 decisions.



GUI controls can report success before mutation truth is known

Channel connect/disconnect paths can advance or toast without confirming all API mutation results; prompt overrides write on every change.






Exhibit 25 · Evidence readout

Sources C9, C10. Decision sequence – User action → API mutation → Confirm result → Success/error state

Management implication

Require confirmed ok, error recovery, and explicit save/cancel boundaries.



Operational history resets and cannot correlate one run end-to-end

Bridge metrics are process-memory, event history is bounded, and message→thread→phase→subagent→receipt→delivery IDs are not one durable chain.






Exhibit 26 · Evidence readout

Sources C12, O4, S7. Decision sequence – Ingress → Codex turn → PABCD → Subagent/approval → Receipt → Delivery

Management implication

Persist local redacted events and expose timeline/doctor/export.



Model/provider controls are rich but vocabulary drifts across surfaces

Subagents and messenger agents expose different effort sets, while capability truth is not one shared contract.






Exhibit 27 · Evidence readout

Sources C18, O9, S8. Decision sequence – Catalog → Capability → Role policy → GUI/CLI/MCP → Spawn

Management implication

One model/effort vocabulary with unavailable states and source of capability truth.



Source CI is broad; real Windows install lifecycle is missing

Windows source tests and WSL are strong, but install/retrust/upgrade/remove runs only on Linux and macOS.

Evidence base

C15, C17, O7, S5

Diagnostic question

Close candidate platform proof.

Decision test

Keep if exact platform gap.

D/E. Platform & scale maturity (1–5; N/E excluded)

CodexClaw

4

OMO

4

Senpi

4

Exhibit 28 · Evidence readout

Sources C15, C17, O7, S5. Directional maturity – CodexClaw 4; OMO 4; Senpi 4. Linux: source + install

Management implication

Add exact-SHA Windows marketplace lifecycle before publication.

Evidence-backed diagnosis

- Linux: source + install
- macOS: source + install

Accountable owner

platform/release

Primary KPI

3/3 OS lifecycle

Principal risk

source CI false confidence

Distribution is candidate-grade except checkout-only map/GUI surfaces

Payload freshness and release proof are strong; installed map/GUI claims remain inconsistent with command availability.






Exhibit 29 · Evidence readout

Sources C14, C17, O10. Decision sequence – Source → Built payload → Marketplace cache → Installed commands → User-visible claims

Management implication

Ship bounded commands or remove installed affordances and documentation claims.



Maintainability is the limiting engineering constraint

No typecheck/lint/coverage gate, 21 plugin-src files over 400 LOC, four over 800, and four win-exec owners raise change risk.

Evidence base

C16, O7, S5

Diagnostic question

Fund the static-safety spine.

Decision test

Keep if top engineering gap.

D. Engineering platform maturity (1–5; N/E excluded)

CodexClaw

2

OMO

3

Senpi

4

Exhibit 30 · Evidence readout

Sources C16, O7, S5. Directional maturity – CodexClaw 2; OMO 3; Senpi 4. Baseline: 21 >400, 4 >800, 4 win-exec

Management implication

Typecheck, modularize, canonicalize helpers, stage CI, then expand runtime surfaces.

Evidence-backed diagnosis

- Baseline: 21 >400, 4 >800, 4 win-exec
- Target: ≤8, 0, 1

Accountable owner

build/CI

Primary KPI

0 diagnostics; ≤8 files >400 LOC; 0 >800; 1 win-exec owner

Principal risk

feature growth before repair

Performance measurement exists but does not govern candidates

Hook benchmark and comparator exist, but no persistent same-runner baseline or release budget makes them operational evidence.

Evidence base C19, S9	Diagnostic question Establish advisory trend before a blocking gate.	Decision test Keep if measured-operations decision.
---------------------------------	--	---

E/D. Scale & platform maturity (1–5; N/E excluded)




Exhibit 31 · Evidence readout

Sources C19, S9. Directional maturity – CodexClaw 2; OMO 3; Senpi 3. Cold/warm p50/p95 tools exist

Management implication

Collect >=30 comparable runs/OS, then enforce per-hook regression budgets.

Evidence-backed diagnosis

- Cold/warm p50/p95 tools exist
- Windows CI benchmark skipped

Accountable owner build/CI	Primary KPI >=30 comparable runs/OS	Principal risk benchmark theatre
--------------------------------------	---	--

Sources: C19, S9 ·

Internal capability breadth exceeds the external ecosystem contract

Twenty-eight internal skills and search/show exist, but external bodies are mutable and no stable extension ABI/trust receipt exists.

Evidence base

C20, O10, S10

Diagnostic question

Define a bounded extension contract.

Decision test

Keep if ecosystem decision.

E. Scale & ecosystem maturity (1–5; N/E excluded)

CodexClaw

3

OMO

4

Senpi

3

Exhibit 32 · Evidence readout

Sources C20, O10, S10. Directional maturity – CodexClaw 3; OMO 4; Senpi 3. Internal breadth: strong

Management implication

Version manifest, provenance, compatibility, disable path, permission diff, and uninstall proof.

Evidence-backed diagnosis

- Internal breadth: strong
- External trust/ABI: partial

Accountable owner

release/security

Primary KPI

immutable skill rate=100%

Principal risk

mutable extension execution

Impact × feasibility favors five programs over feature parity

Unlike page 3's executive thumbnail, this exhibit tests sequencing and feasibility against the diagnosed domains.

Evidence base

C2, C10, C12, C16, C17, C19, C20

Diagnostic question

Approve portfolio sequencing.

Decision test

Keep if portfolio choice.

Impact ↑




Exhibit 33 · Evidence readout

Sources C2, C10, C12, C16, C17, C19, C20. Decision sequence – P1 Truth → P2 UX → P3 Ops → P4 Safety → P5 Delivery

Management implication

Prioritize P4/P1 foundations, then P2/P3 product integration, with P5 candidate gates in parallel.

Accountable owner

transformation lead + program owners

Primary KPI

five program exit gates, reviewed monthly

Principal risk

portfolio dilution or runtime duplication

0–90 days should close truth, static safety, mutation UX, and map contradiction

Near-term work reduces risk and closes concrete product inconsistencies before broader state or console growth.

Workstream	0–30 days	31–60 days	61–90 days	Exit evidence
P4 Static safety	Typecheck + lint baseline	Module debt owners	CI enforcement	0 diagnostics; debt trend
P1 Truth	Capability snapshot	Public goalplan ops	Bound review policy	0 misroutes; lifecycle API
P2 UX	Mutation contract	False-success tests	Healthy-turn funnel	0 false success; <5m
P5 Delivery	Map contradiction decision	Windows install smoke	3-OS lifecycle	3/3 OS; docs aligned

Exhibit 34 · Evidence readout

Sources C2, C9, C14, C16, C17. P4 static safety, P1 truth, P2 UX, and P5 delivery each carry 30/60/90-day outputs and an evidence-based exit gate.

Management implication

Start static checks and capability truth first; ship public goalplan ops and correct mutation/map UX by day 90.

Accountable owner

transformation lead + program owners

Primary KPI

five program exit gates, reviewed monthly

Principal risk

portfolio dilution or runtime duplication

3–6 months should build durable operations, operator console, and security release

Scale initiatives depend on stable IDs, structured state, privacy model, and static-safety foundations.

Workstream	Month 3–4	Month 5	Month 6	Exit evidence
P3 Operations	Stable correlation IDs	Full-chain trace	SLO review	>=99% traced/completed
P2 Console	Read-only state projection	Operator actions	Recovery UX	<2s state; 0 false success
P5 Security	SBOM + provenance	Signing + verify	Candidate gate	100% signed assets
P5 Ecosystem	Manifest + permissions	Compatibility + disable	Bounded pilot	100% immutable approved

Exhibit 35 · Evidence readout

Sources C6, C7, C10, C12. P3 operations precedes P2 console; P5 signing precedes ecosystem pilot. Month-six gates are >=99% trace, <2s state, and 100% signed assets.

Management implication

Build correlated operations before the unified console; add SBOM/signing before ecosystem expansion.

Accountable owner

transformation lead + program owners

Primary KPI

five program exit gates, reviewed monthly

Principal risk

portfolio dilution or runtime duplication

An operating model ties each program to owner, cadence, formula, and gate

Five programs have accountable owners, reproducible formulas, data sources, review cadence, and entry/exit criteria.

Program	Accountable owner	Metric formula and window	Review cadence	Exit gate
P1 Truth	pabcd-state	misrouted governed transitions / all governed transitions; rolling 30 days	Per PR + weekly	0 misroutes; public add/resolve/capture lifecycle
P2 UX	bridge + GUI	median minutes from install start to first healthy turn; release-candidate cohort	Weekly + each RC	<5 minutes; 0 false-success mutations
P3 Ops	bridge operations	fully correlated terminal outcomes / completed remote turns; rolling 7 days	Daily SLO + weekly	>=99% traced; every failure classified
P4 Safety	build + CI	compiler/lint diagnostics; source files >400/>800 LOC; win-exec owners; monthly snapshot	Per PR + monthly	0 diagnostics; <=8/>400, 0/>800, 1 owner
P5 Delivery	release + security	OS lifecycle passes / 3; signed assets / published assets; each candidate	Each candidate + monthly	3/3 OS; 100% signed; SBOM attached

Exhibit 36 · Evidence readout

Sources C2, C10, C12, C16, C17. P1 uses 30-day misroutes/governed transitions; P2 uses cohort median install-to-healthy-turn; P3 uses seven-day traced/completed turns.

Management implication

Manage outcomes through formulas and decision gates, not feature completion narratives.

Accountable owner

transformation lead + program owners

Primary KPI

five program exit gates, reviewed monthly

Principal risk

portfolio dilution or runtime duplication

Selective transfer preserves CodexClaw's differentiation

Adopt artifacts and information models; adapt state and evidence patterns; defer host-dependent mechanisms; reject duplicate runtime ownership.

ADOPT	ADAPT	DEFER	REJECT
Executable plans	Claim lineage	Named deep arm	Duplicate scheduler
QA taxonomy	Operator info model	Team DAG	Second goal/task DB
Tool-aware composition	Perf trends	Activation telemetry	Global parser/arming

Exhibit 37 · Evidence readout

Sources C1, O1, O3, O4, S4, S7. Adopt executable plans and QA taxonomy; adapt claim lineage and operator models; defer host-dependent modes; reject duplicate runtime owners.

Management implication

Import patterns, not scheduler/goal/terminal/TUI owners.

Accountable owner

transformation lead + program owners

Primary KPI

five program exit gates, reviewed monthly

Principal risk

portfolio dilution or runtime duplication

Sources: C1, O1, O3, O4, S4, S7 ·

Risks and bounded uncertainty constrain the recommendation

Competitor evidence is asymmetric, current HEAD is not remote-certified, and several runtime/test reports remain source-traced rather than live-proven.

Evidence base

C5, C6, C7, O7, S5

Diagnostic question

Accept risks and evidence limits.

Decision test

Keep if risk acceptance.

Impact ↑




Exhibit 38 · Evidence readout

Sources C5, C6, C7, O7, S5. Decision sequence – Evidence asymmetry → Current HEAD uncertified → Host-private liveness → Supply-chain expansion → Ceremony risk

Management implication

Keep N/E visible, require candidate receipts, and avoid extrapolating peer runtime behavior.

Accountable owner

transformation lead + program owners

Primary KPI

five program exit gates, reviewed monthly

Principal risk

portfolio dilution or runtime duplication

Appendix A – Evidence method and score aggregation

Methodology fixes the comparison boundary and prevents prose, absent evidence, or asymmetric competitor scope from becoming numeric certainty.

Stage	Definition	Use
1	Absent or prompt-only	Do not claim runtime
2	Partial/manual/test-only	Gap remains operational
3	Shipped bounded runtime	Production path exists
4	Durable integrated evidence	Cross-surface owner and proof
5	Measured operating system	Persistent feedback changes decisions
N/E	Evidence insufficient	Excluded, never zero

Aggregation: N/E is excluded. No product-wide average is published. CodexClaw pillar averages use only assessed domains and are descriptive, not decision gates.

Evidence hierarchy: current shipped source + reachable test/runtime path > current source with incomplete reachability > historical/asymmetric evidence. Prompt-only is never promoted to runtime.

Code	Source
C1	plugins/codexclaw/components/pabcd-state/src/fsm.ts:23-70; hook.ts:1275-1336
C2	plugins/codexclaw/components/pabcd-state/src/goalplan.ts:59-90,700-749; goalplan-cli.ts:37-84
C3	plugins/codexclaw/skills/search/SKILL.md:118-177; components/recall/src/chat-search.ts:28-97
C4	plugins/codexclaw/components/subagent-config/src/store.ts:16-58; spawn-attach-hook.ts:745-929
C5	plugins/codexclaw/components/pabcd-state/src/source-identity.ts:168-208; goalplan.ts:921-994
C6	plugins/codexclaw/components/cxc-ops/src/hook-trust.ts:308-467; docs/security-hardening.md:86-95
C7	.github/workflows/release.yml:68-225; packed-install.yml:21-160
C8	plugins/codexclaw/components/pabcd-state/src/interview.ts:264-340; goal-active.ts:57-89
C9	README.md:48-61; plugins/codexclaw/gui/src/pages/Channels.tsx:216-295
C10	plugins/codexclaw/gui/src/App.tsx:20-25; pages/Dashboard.tsx:181-243
C11	messenger-bridge/src/runner.ts:78-105,542-573; test/runner.test.ts:37-53,166-184; src/gateway-commands.ts:50-66; test/gateway-commands.test.ts:251-338
C12	plugins/codexclaw/components/messenger-bridge/src/event-log.ts:27-45,112-139; metrics.ts:1-20,42-90
C13	plugins/codexclaw/components/subagent-config/src/capabilities.ts:54-128; spawn-wrapper.ts:428-500
C14	plugins/codexclaw/components/cxc-ops/src/map-affordance.ts:101-120; plugins/codexclaw/bin/cxc.mjs:89-139

Appendix B – Every management page resolves to evidence codes

Pages 2–39 map decisions to exact source families; pages without distinct decisions were merged before build.

Pg	Evidence
2	C1,C5,C17,O4,S7
3	C2,C10,C12,C16,C17
4	C5,C16,O7,S9
5	C1,C5,C10,C16,C17
6	C1,C4,O4,S7
7	C1,C5,C10,C12,C16,C19
8	C2,C5,C10,C12
9	C1,O1
10	C2,O2,S1
11	C1,C5,O5
12	C3,O3
13	C3,C13,S6
14	C3,O3
15	C4,C20,O11
16	C14
17	C4,C13,O4
18	C12,O4,S7
19	C5,O5
20	C5,O5

Pg	Evidence
21	C6,S4
22	C6,C7,S4
23	C11
24	C10,C2,O8
25	C9,O8,S10
26	C9,C10
27	C12,O4,S7
28	C18,O9,S8
29	C15,C17,O7,S5
30	C14,C17,O10
31	C16,O7,S5
32	C19,S9
33	C20,O10,S10
34	C2,C10,C12,C16,C17,C19,C20
35	C2,C9,C14,C16,C17
36	C6,C7,C10,C12
37	C2,C10,C12,C16,C17
38	C1,O1,O3,O4,S4,S7
39	C5,C6,C7,O7,S5

Code	Source
C15	.github/workflows/ci.yml:12-55; wsl.yml:13-40
C16	package.json:21-24; plugins/codexclaw/components/pabcd-state/src/hook.ts:1469; goalplan.ts:1094
C17	.agents/plugins/marketplace.json:6-18; plugins/codexclaw/test/payload-bin.test.mjs:53-106
C18	plugins/codexclaw/components/provider-bridge/src/detect.ts:1-101; plugins/codexclaw/gui/src/api.ts:12-16,239-241
C19	plugins/codexclaw/scripts/hook-bench.mjs:123-160; hook-bench-compare.mjs:5-40
C20	plugins/codexclaw/components/skill-search/src/cli.ts:153-230; sources.ts:10-20
O1	devlog/omo/packages/omo-codex/plugin/components/ulw-loop/src/checkpoint.ts:158-252
O2	devlog/omo/packages/team-core/src/team-tasklist/claim.ts:45-96; omo-opencode/src/tools/task/types.ts:3-21
O3	devlog/omo/packages/shared-skills/skills/ulw-research/SKILL.md:61-149,195-260
O4	devlog/omo/packages/omo-opencode/src/features/background-agent/manager.ts:245-320,587-655
O5	devlog/omo/packages/omo-codex/plugin/components/ulw-loop/src/quality-gate.ts:93-205
O7	devlog/omo/.github/workflows/ci.yml:145-378; devlog/omo/package.json:40-47,224-250
O8	devlog/omo/packages/omo-opencode/src/cli/tui-installer.ts:22-210; src/tui.ts:118-184
O9	devlog/omo/packages/omo-opencode/src/config/schema/agent-overrides.ts:6-67; categories.ts:5-40
O10	devlog/omo/package.json:8-40; packages/omo-codex/package.json:2-32; packages/omo-senpi/package.json:2-56
O11	devlog/omo/packages/skills-loader-core/src/features/opencode-skill-loader/async-loader.ts:74-100; omo-opencode/src/tools/delegate-task/skill-resolver.ts:71-114

Appendix C – Detailed score ledger and bounded uncertainty

The ledger retains N/E and confidence; exact OMO-pinned Senpi integration, current remote certification, and controlled-task performance remain outside proven scope.

Domain	CXC	OMO	Senpi	Conf.	Codes
Orchestration	4	4	3	H/H/M	C1,O1,S1
Goal/task	3	4	4	H/H/H	C2,O2,S1
Research/knowledge	3	4	2	H/H/M	C3,O3,S2
Multi-agent	3	4	3	M/H/M	C4,O4,S3
Evidence/QA	4	4	3	H/H/M	C5,O5,S1
Trust/security	3	3	4	H/M/H	C6,O5,S4
Release	4	4	4	H/M/H	C7,O7,S5
Authority	4	3	4	H/M/H	C8,O1,S4
Onboarding	3	4	N/E	M/H/-	C9,O8
Console	2	4	N/E	M/H/-	C10,O8
Remote	4	N/E	N/E	H/-/-	C11
Observability	2	4	4	H/H/H	C12,O4,S7
Routing	2	4	4	M/H/H	C13,O11,S6
Workspace intel	2	N/E	N/E	H/-/-	C14
Cross-platform	4	4	4	H/H/H	C15,O7,S5
Maintainability	2	3	4	H/M/H	C16,O7,S5
Distribution	4	4	4	H/H/H	C17,O10,S5
Model/provider	3	4	4	M/H/H	C18,O9,S8
Performance	2	3	3	H/M/H	C19,O7,S9
Ecosystem	3	4	3	M/H/M	C20,O10,S10

Bounded uncertainty: exact OMO peer Senpi 2026.8.26-2 was not materialized; remote GitHub exact-head state is not refreshed here; competitor maintainability and remote-operations evidence is asymmetric; model compliance and host implicit skill loading remain unobservable.

Code	Source
S1	devlog/.senpi/packages/coding-agent/src/core/extensions/builtin/goal/store.ts:19-82,148-208; todo-gate.ts:8-30
S2	devlog/.senpi/packages/coding-agent/src/core/dynamic-prompt/intent-gate.ts:14-45; parallel-tools.ts:1-6
S3	devlog/.senpi/packages/agent/src/agent-loop.ts:846-993; devlog/.omo/packages/senpi-task/src/lifecycle/shutdown.ts:9-67
S4	devlog/.senpi/packages/coding-agent/src/core/extensions/builtin/permission-system/index.ts:102-155; core/project-trust.ts:46-95
S5	devlog/.senpi/github/workflows/ci.yml:13-187; build-binaries.yml:81-140
S6	devlog/.senpi/packages/coding-agent/src/core/dynamic-prompt/build.ts:61-101; tool-section.ts:12-45
S7	devlog/.senpi/packages/coding-agent/src/core/extensions/builtin/terminal/monitor-registry.ts:55-164; goal/monitor-continuation.ts:323-618
S8	devlog/.senpi/package.json:15-45; packages/coding-agent/package.json:10-50
S9	devlog/.senpi/github/workflows/perf-trend.yml:1-13,82-88
S10	devlog/.senpi/packages/coding-agent/src/core/extensions/builtin/index.ts:61-100; devlog/.senpi/package.json:5-14