

CodexClaw 전체 격차 전략 보고서

Governance-first thin harness를 유지하면서 5개 전략 프로그램으로
제품·운영·엔지니어링 격차를 닫는다.

기준	고정 범위
CodexClaw	98cac96 · local current source
OMO beta	84f98d8 · v5.0.0-beta.22
Senpi	703d9d7 · 2026.8.27 (OMO peer 2026.8.26-2)
범위	5 pillars · 20 domains · 42 pages

2026-08-27 · Directional maturity assessment, not a performance benchmark

CodexClaw should scale governance, not recreate the runtime

The product's advantage is durable policy and evidence over native Codex; the next stage is integration, operator experience, and measured delivery.




Exhibit 1. Text alternative: Set the strategic thesis. Invest in control-plane truth and product operations; reject duplicate runtime ownership.

Management implication

Invest in control-plane truth and product operations; reject duplicate runtime ownership.

- Strongest: evidence and bounded orchestration
- Weakest: connected state, UX, observability
- Boundary: host keeps scheduler/terminal/goal authority

Five programs close the largest gaps without breaking the thin-host boundary

The portfolio addresses truth, experience, operations, engineering safety, and candidate-grade delivery.




Exhibit 2. Text alternative: Approve the five-program portfolio. Fund five programs in dependency order rather than pursuing competitor feature parity.

Management implication

Fund five programs in dependency order rather than pursuing competitor feature parity.

Evidence supports directional maturity, not benchmark rankings

Scores classify evidence maturity from prose to measured operations; N/E is excluded rather than scored zero.




Exhibit 3. Text alternative: Bound how scores may be used. Use the heatmap to prioritize domains, not to publish a product winner.

Management implication

Use the heatmap to prioritize domains, not to publish a product winner.

Five MECE pillars define what 'full gap' means

Twenty domains cover strategy, assurance, product operations, engineering platform, and scale/ecosystem without double-counting owners.




Exhibit 4. Text alternative: Approve completeness denominator. A page only survives if it changes priority, owner, KPI, transfer, or gate.

Management implication

A page only survives if it changes priority, owner, KPI, transfer, or gate.

CodexClaw owns policy and evidence; Codex owns execution

The owner boundary is the central strategic differentiator and the filter for every transfer decision.



Exhibit 5. Text alternative: Preserve the owner boundary. Do not create a second scheduler, terminal registry, goal DB, or Stop owner.

Management implication

Do not create a second scheduler, terminal registry, goal DB, or Stop owner.

Maturity is bimodal: governance is strong, product operations are fragmented

CodexClaw scores 4 in orchestration/evidence/release, but 2 in console, observability, routing, workspace intelligence, maintainability, and performance operations.

	오케스트레이션	연구	증거	UX	관측	배포	유지보수	생태계
CodexClaw	4	3	4	3	2	4	2	3
OMO	4	4	4	4	4	4	3	4
Senpi	3	2	3	3	4	4	4	4

Exhibit 6. Text alternative: Identify portfolio imbalance. Balance the control plane with product operations and engineering discipline.

Management implication

Balance the control plane with product operations and engineering discipline.

CodexClaw wins where truth is durable and loses where state is disconnected

The causal gap is not missing features; it is weak connection between existing PABCD, bridge, capability, evidence, and delivery state.




Exhibit 7. Text alternative: Prioritize integration over surface growth. Connect authoritative state before adding new modes or agents.

Management implication

Connect authoritative state before adding new modes or agents.

PABCD is real runtime, but intellectual quality remains partly prompt-bound

Legal transitions and receipts are mechanized; dependency slicing, verifier realism, and adversarial reasoning still depend on model compliance.

A. Strategy & execution maturity (1–5; N/E excluded)



Exhibit 8. Text alternative: Tier governance by work risk. Keep lean flow for C0–C2; require bound review and stronger evidence for C3/C4.

Management implication

Keep lean flow for C0–C2; require bound review and stronger evidence for C3/C4.

- Shipped: FSM, plan gate, check gate
- Prompt: plan quality, synthesis, reviewer reuse

Goalplan blocks false completion but cannot operate the full task lifecycle

Tasks and criteria affect completion, yet public add-task, resolve-task, and evidence-capture operations do not exist.




Exhibit 9. Text alternative: Fund public task and criterion operations. Extend the existing authoritative goalplan; do not add a second task database.

Management implication

Extend the existing authoritative goalplan; do not add a second task database.

- Current: pending/done checklist
- Needed: owner, outcome, evidence, retry lineage

Review identity is strong when bound and form-only when lean

Launch IDs, plan hashes, and reviewer identity are durable only when a review round is opened; lean attestation remains unauthenticated text.




Exhibit 10. Text alternative: Define a tiered review policy. Require current-plan review binding for C3/C4 while preserving bounded recovery.

Management implication

Require current-plan review binding for C3/C4 while preserving bounded recovery.

Research quality rules outpace the state that must recover them

Tier 3 specifies waves, leads, and claims, but compaction cannot deterministically recover the next research action.

A. Strategy & execution maturity (1–5; N/E excluded)




Exhibit 11. Text alternative: Fund durable research projection. Map run→work phase, lane/lead→task, claim→criterion, ledger→history.

Management implication

Map run→work phase, lane/lead→task, claim→criterion, ledger→history.

- CodexClaw: proof ladder, no durable wave state
- OMO: deep claim/observation lineage

Search proof is strong; capability-aware routing is weak

Discovery-versus-proof doctrine is excellent, but the live search ladder is not rewritten from the tools actually available this turn.




Exhibit 12. Text alternative: Connect runtime capability truth to search. Maintain an explicit unknown state and reject silent tool fallback.

Management implication

Maintain an explicit unknown state and reject silent tool fallback.

Recall is mature retrieval, not epistemic truth maintenance

FTS, memory ranking, and compaction reinjection are shipped; durable claims still lack validity, counterevidence, and corroboration fields.

A. Strategy & execution maturity (1–5; N/E excluded)



Exhibit 13. Text alternative: Add claim validity before memory promotion. Promote only proven claims with source validity; keep retrieval and truth separate.

Management implication

Promote only proven claims with source validity; keep retrieval and truth separate.

- Strength: CWD-scoped untrusted recall
- Gap: stale curated memory can outrank newer evidence

Skill delivery is robust; external skill trust remains mutable

Mentioned built-in skills are bounded and inlined; external discovery fetches mutable bodies without immutable revision or permission receipt.




Exhibit 14. Text alternative: Harden skill supply-chain review. Pin revision/hash/license/permissions and separate review from activation.

Management implication

Pin revision/hash/license/permissions and separate review from activation.

Workspace intelligence is advertised but incomplete in the installed payload

SessionStart and plugin prompts advertise cxc map, while the payload dispatcher refuses the command as checkout-only.




Exhibit 15. Text alternative: Decide ship or stop advertising. Either ship a bounded payload map smoke or remove installed affordances.

Management implication

Either ship a bounded payload map smoke or remove installed affordances.

Multi-agent policy is strong; runtime-surface detection can misroute

Role stores, leaf guards, and evidence gates are shipped; config-only V1/V2 inference and timeout semantics can choose the wrong lifecycle path.

A/D. Execution & platform maturity (1–5; N/E excluded)



Exhibit 16. Text alternative: Build a session capability adapter. Snapshot effective schema, identity, lifecycle verbs, concurrency, and unavailable capabilities before spawn.

Management implication

Snapshot effective schema, identity, lifecycle verbs, concurrency, and unavailable capabilities before spawn.

- Strength: bounded roles + skill transport
- Gap: surface truth and wait semantics

Host-native background execution lacks an integrated liveness projection

Codex owns workers and terminals, but CodexClaw lacks an authenticated cross-surface view of live child/terminal status.




Exhibit 17. Text alternative: Keep scheduling host-owned and request live events. Consume authenticated host snapshots; never synthesize a local wake bus.

Management implication

Consume authenticated host snapshots; never synthesize a local wake bus.

Source provenance leads, while receipt truth is still self-authored

Exact source identity invalidates stale evidence, but a writer can still claim a command and result that were never independently observed.




Exhibit 18. Text alternative: Build Evidence Authority v3. Bind command digest, output digest, exit status, executor identity, timing, and review round.

Management implication

Bind command digest, output digest, exit status, executor identity, timing, and review round.

Final gates are sophisticated, but the strongest schema remains optional

Schema-v2 binds reviewer, test, QA, and source; marker deletion and fail-open paths keep it from being universal.

B. Assurance & trust maturity (1–5; N/E excluded)



Exhibit 19. Text alternative: Mandate the strongest gate by work class. Require schema v2 for C3/C4 and release work; keep bounded escape paths visible as non-success.

Management implication

Require schema v2 for C3/C4 and release work; keep bounded escape paths visible as non-success.

- C0–C2: lean
- C3–C4: source-bound
- Release: exact-head + publish proof

Hook trust protects local integrity, not publisher authenticity

Hash trust and atomic retrust prove installed consistency; they do not establish a signed publisher-to-payload chain.




Exhibit 20. Text alternative: Add supply-chain identity. Bind trusted hooks to SBOM, signature, release SHA, and post-install verification.

Management implication

Bind trusted hooks to SBOM, signature, release SHA, and post-install verification.

Security release assurance lags the repository's stated standard

CI/release includes tests, inventory, and packaging but lacks dedicated SAST, dependency audit, secret scan, SBOM, and signing.

B. Assurance & trust maturity (1–5; N/E excluded)



Exhibit 21. Text alternative: Expand the candidate security gate. Require zero unresolved high/critical findings and signed, verified assets.

Management implication

Require zero unresolved high/critical findings and signed, verified assets.

- Current: exact candidate design
- Missing: SAST/SBOM/signing/secret/dependency proof

Remote operations is a differentiated shipped strength

Telegram/Discord pairing, commands, approval relay, progress, media, sessions, and dashboard are real runtime surfaces.

C. Product & operations maturity (1–5; N/E excluded)

CodexClaw

4

OMO

N/E

Senpi

N/E

Exhibit 22. Text alternative: Preserve and extend a unique advantage. Use the bridge as the entry point for the unified operator console.

Management implication

Use the bridge as the entry point for the unified operator console.

- Shipped: pairing, approvals, progress, sessions
- N/E: no comparable bounded peer surface

The operator console omits the core PABCD product state

The dashboard is bridge-centric and does not surface phase, work phase, criteria, evidence, doctor, or hook trust.




Exhibit 23. Text alternative: Fund a unified console. Show workflow truth and next action without creating parallel state.

Management implication

Show workflow truth and next action without creating parallel state.

Installation is concise but not guided end-to-end

Two commands install the plugin, but hook trust, doctor, model choices, console, and optional bridge are not one progressive journey.




Exhibit 24. Text alternative: Build capability-aware onboarding. Health→trust→models→optional bridge→first successful task in ≤4 decisions.

Management implication

Health→trust→models→optional bridge→first successful task in ≤4 decisions.

GUI controls can report success before mutation truth is known

Channel connect/disconnect paths can advance or toast without confirming all API mutation results; prompt overrides write on every change.




Exhibit 25. Text alternative: Fix control-plane correctness before adding features. Require confirmed ok, error recovery, and explicit save/cancel boundaries.

Management implication

Require confirmed ok, error recovery, and explicit save/cancel boundaries.

Operational history resets and cannot correlate one run end-to-end

Bridge metrics are process-memory, event history is bounded, and message→thread→phase→subagent→receipt→delivery IDs are not one durable chain.




Exhibit 26. Text alternative: Fund a durable correlated operations plane. Persist local redacted events and expose timeline/doctor/export.

Management implication

Persist local redacted events and expose timeline/doctor/export.

Model/provider controls are rich but vocabulary drifts across surfaces

Subagents and messenger agents expose different effort sets, while capability truth is not one shared contract.




Exhibit 27. Text alternative: Unify capability-aware settings. One model/effort vocabulary with unavailable states and source of capability truth.

Management implication

One model/effort vocabulary with unavailable states and source of capability truth.

Source CI is broad; real Windows install lifecycle is missing

Windows source tests and WSL are strong, but install/retrust/upgrade/remove runs only on Linux and macOS.

D/E. Platform & scale maturity (1–5; N/E excluded)



Exhibit 28. Text alternative: Close candidate platform proof. Add exact-SHA Windows marketplace lifecycle before publication.

Management implication

Add exact-SHA Windows marketplace lifecycle before publication.

- Linux: source + install
- macOS: source + install
- Windows: source only

Distribution is candidate-grade except checkout-only map/GUI surfaces

Payload freshness and release proof are strong; installed map/GUI claims remain inconsistent with command availability.




Exhibit 29. Text alternative: Decide the shipped product surface. Ship bounded commands or remove installed affordances and documentation claims.

Management implication

Ship bounded commands or remove installed affordances and documentation claims.

Maintainability is the limiting engineering constraint

No typecheck/lint/coverage gate, 21 plugin-src files over 400 LOC, four over 800, and four win-exec owners raise change risk.

D. Engineering platform maturity (1–5; N/E excluded)




Exhibit 30. Text alternative: Fund the static-safety spine. Typecheck, modularize, canonicalize helpers, stage CI, then expand runtime surfaces.

Management implication

Typecheck, modularize, canonicalize helpers, stage CI, then expand runtime surfaces.

- Baseline: 21 >400, 4 >800, 4 win-exec
- Target: ≤8, 0, 1

Performance measurement exists but does not govern candidates

Hook benchmark and comparator exist, but no persistent same-runner baseline or release budget makes them operational evidence.

E/D. Scale & platform maturity (1–5; N/E excluded)




Exhibit 31. Text alternative: Establish advisory trend before a blocking gate. Collect ≥ 30 comparable runs/OS, then enforce per-hook regression budgets.

Management implication

Collect ≥ 30 comparable runs/OS, then enforce per-hook regression budgets.

- Cold/warm p50/p95 tools exist
- Windows CI benchmark skipped
- No candidate gate

Internal capability breadth exceeds the external ecosystem contract

Twenty-eight internal skills and search/show exist, but external bodies are mutable and no stable extension ABI/trust receipt exists.

E. Scale & ecosystem maturity (1–5; N/E excluded)



Exhibit 32. Text alternative: Define a bounded extension contract. Version manifest, provenance, compatibility, disable path, permission diff, and uninstall proof.


Management implication

Version manifest, provenance, compatibility, disable path, permission diff, and uninstall proof.

- Internal breadth: strong
- External trust/ABI: partial

Impact × feasibility favors five programs over feature parity

The highest-value initiatives reuse existing owners; competitor runtime parity scores lower feasibility and higher architectural risk.



Management implication

Prioritize P4/P1 foundations, then P2/P3 product integration, with P5 candidate gates in parallel.

0–90 days should close truth, static safety, mutation UX, and map contradiction

Near-term work reduces risk and closes concrete product inconsistencies before broader state or console growth.

0–30 static checks + truth

61–90 Windows lifecycle + map decision



Exhibit 34. Text alternative: Approve near-term roadmap. Start static checks and capability truth first; ship public goalplan ops and correct mutation/map UX by day 90.

Management implication

Start static checks and capability truth first; ship public goalplan ops and correct mutation/map UX by day 90.

3–6 months should build durable operations, operator console, and security release

Scale initiatives depend on stable IDs, structured state, privacy model, and static-safety foundations.




Exhibit 35. Text alternative: Approve scale roadmap. Build correlated operations before the unified console; add SBOM/signing before ecosystem expansion.

Management implication

Build correlated operations before the unified console; add SBOM/signing before ecosystem expansion.

An operating model ties each program to owner, cadence, formula, and gate

Five programs have accountable owners, measurable baselines, data sources, review cadence, and entry/exit criteria.

Program	Owner	Cadence	Exit gate
P1 Truth	pabcd-state	PR/weekly	zero misroutes; public lifecycle
P2 UX	bridge/gui	weekly/RC	<5m first turn; zero false success
P3 Ops	bridge ops	daily/weekly	>=99% correlated trace
P4 Safety	build/CI	PR/monthly	0 diagnostics; debt targets
P5 Delivery	release/security	candidate/monthly	3/3 OS; signed assets

Exhibit 36. Text alternative: Approve governance and cadence. Manage outcomes through formulas and decision gates, not feature completion narratives.

Management implication

Manage outcomes through formulas and decision gates, not feature completion narratives.

Selective transfer preserves CodexClaw's differentiation

Adopt artifacts and information models; adapt state and evidence patterns; defer host-dependent mechanisms; reject duplicate runtime ownership.

ADOPT	ADAPT	DEFER	REJECT
Executable plans	Claim lineage	Named deep arm	Duplicate scheduler
QA taxonomy	Operator info model	Team DAG	Second goal/task DB
Tool-aware composition	Perf trends	Activation telemetry	Global parser/arming

Exhibit 37. Text alternative: Approve transfer decisions. Import patterns, not scheduler/goal/terminal/TUI owners.

Management implication
Import patterns, not scheduler/goal/terminal/TUI owners.

Risks and bounded uncertainty constrain the recommendation

Competitor evidence is asymmetric, current HEAD is not remote-certified, and several runtime/test reports remain source-traced rather than live-proven.




Exhibit 38. Text alternative: Accept risks and evidence limits. Keep N/E visible, require candidate receipts, and avoid extrapolating peer runtime behavior.

Management implication

Keep N/E visible, require candidate receipts, and avoid extrapolating peer runtime behavior.

Appendix A – Evidence method and score aggregation

Methodology fixes the comparison boundary and prevents prose, absent evidence, or asymmetric competitor scope from becoming numeric certainty.

Stage	Definition	Use
1	Absent or prompt-only	Do not claim runtime
2	Partial/manual/test-only	Gap remains operational
3	Shipped bounded runtime	Production path exists
4	Durable integrated evidence	Cross-surface owner and proof
5	Measured operating system	Persistent feedback changes decisions
N/E	Evidence insufficient	Excluded, never zero

Aggregation: N/E is excluded. No product-wide average is published. CodexClaw pillar averages use only assessed domains and are descriptive, not decision gates.

Evidence hierarchy: current shipped source + reachable test/runtime path > current source with incomplete reachability > historical/asymmetric evidence. Prompt-only is never promoted to runtime.

Code	Source
C1	plugins/codexclaw/components/pabcd-state/src/fsm.ts:23-70; hook.ts:1275-1336
C2	plugins/codexclaw/components/pabcd-state/src/goalplan.ts:59-90,700-749; goalplan-cli.ts:37-84
C3	plugins/codexclaw/skills/search/SKILL.md:118-177; components/recall/src/chat-search.ts:28-97
C4	plugins/codexclaw/components/subagent-config/src/store.ts:16-58; spawn-attach-hook.ts:745-929
C5	plugins/codexclaw/components/pabcd-state/src/source-identity.ts:168-208; goalplan.ts:921-994
C6	plugins/codexclaw/components/cxc-ops/src/hook-trust.ts:308-467; docs/security-hardening.md:86-95
C7	.github/workflows/release.yml:68-225; packed-install.yml:21-160
C8	plugins/codexclaw/components/pabcd-state/src/interview.ts:264-340; goal-active.ts:57-89
C9	README.md:48-61; plugins/codexclaw/gui/src/pages/Channels.tsx:216-295
C10	plugins/codexclaw/gui/src/App.tsx:20-25; pages/Dashboard.tsx:181-243
C11	plugins/codexclaw/components/messenger-bridge/README.md:1-30,73-107
C12	plugins/codexclaw/components/messenger-bridge/src/event-log.ts:27-45,112-139; metrics.ts:1-20,42-90

Appendix B – Every management page resolves to evidence codes

Pages 2–39 map decisions to exact source families; pages without distinct decisions were merged before build.

Pg	Evidence
2	C1,C5,C17,O4,S7
3	004
4	003,C5
5	000
6	C1,C4,O4,S7
7	003
8	C2,C5,C10,C12
9	C1,O1
10	C2,O2,S1
11	C1,C5,O5
12	C3,O3
13	C3,C13,S6
14	C3,O3
15	C4,C20,O11
16	C14
17	C4,C13,O4
18	C12,O4,S7
19	C5,O5
20	C5,O5

Pg	Evidence
21	C6,S4
22	C6,C7,S4
23	C11
24	C10,C2,O8
25	C9,O8,S10
26	C9,C10
27	C12,O4,S7
28	C18,O9,S8
29	C15,C17,O7,S5
30	C14,C17,O10
31	C16,O7,S5
32	C19,S9
33	C20,O10,S10
34	004,pages 9–33
35	004
36	004
37	004
38	C1,O1,O3,O4,S4,S7
39	003,UNVERIFIED

Code	Source
C13	plugins/codexclaw/components/subagent-config/src/capabilities.ts:54-128; spawn-wrapper.ts:428-500
C14	plugins/codexclaw/components/cxc-ops/src/map-affordance.ts:101-120; plugins/codexclaw/bin/cxc.mjs:89-139
C15	.github/workflows/ci.yml:12-55; wsl.yml:13-40
C16	package.json:21-24; plugins/codexclaw/components/pabcd-state/src/hook.ts:1469; goalplan.ts:1094
C17	.agents/plugins/marketplace.json:6-18; plugins/codexclaw/test/payload-bin.test.mjs:53-106
C18	plugins/codexclaw/components/provider-bridge/src/detect.ts:1-101; plugins/codexclaw/gui/src/api.ts:12-16,239-241
C19	plugins/codexclaw/scripts/hook-bench.mjs:123-160; hook-bench-compare.mjs:5-40
C20	plugins/codexclaw/components/skill-search/src/cli.ts:153-230; sources.ts:10-20
O1	devlog/omo/packages/omo-codex/plugin/components/ulw-loop/src/checkpoint.ts:158-252
O2	devlog/omo/packages/team-core/src/team-tasklist/claim.ts:45-96; omo-opencode/src/tools/task/types.ts:3-21
O3	devlog/omo/packages/shared-skills/skills/ulw-research/SKILL.md:61-149,195-260
O4	devlog/omo/packages/omo-opencode/src/features/background-agent/manager.ts:245-320,587-655
O5	devlog/omo/packages/omo-codex/plugin/components/ulw-loop/src/quality-gate.ts:93-205
O7	devlog/omo/.github/workflows/ci.yml:145-378; devlog/omo/package.json:40-47,224-250
O8	devlog/omo/packages/omo-opencode/src/cli/tui-installer.ts:22-210; src/tui.ts:118-184
O9	devlog/omo/packages/omo-opencode/src/config/schema/agent-overrides.ts:6-67; categories.ts:5-40
O10	devlog/omo/package.json:8-40; packages/omo-codex/package.json:2-32; packages/omo-senpi/package.json:2-56
O11	devlog/omo/packages/skills-loader-core/src/features/opencode-skill-loader/async-loader.ts:74-100; omo-opencode/src/tools/delegate-task/skill-resolver.ts:71-114

Appendix C – Detailed score ledger and bounded uncertainty

The ledger retains N/E and confidence; exact OMO-pinned Senpi integration, current remote certification, and controlled-task performance remain outside proven scope.

Domain	CXC	OMO	Senpi	Conf.	Codes
Orchestration	4	4	3	H/H/M	C1,O1,S1
Goal/task	3	4	4	H/H/H	C2,O2,S1
Research/knowledge	3	4	2	H/H/M	C3,O3,S2
Multi-agent	3	4	3	M/H/M	C4,O4,S3
Evidence/QA	4	4	3	H/H/M	C5,O5,S1
Trust/security	3	3	4	H/M/H	C6,O5,S4
Release	4	4	4	H/M/H	C7,O7,S5
Authority	4	3	4	H/M/H	C8,O1,S4
Onboarding	3	4	3	H/H/L	C9,O8,S10
Console	2	4	3	H/H/L	C10,O8,S7
Remote	4	N/E	N/E	H/-/-	C11
Observability	2	4	4	H/H/H	C12,O4,S7
Routing	2	4	4	M/H/H	C13,O11,S6
Workspace intel	2	N/E	N/E	H/-/-	C14
Cross-platform	4	4	4	H/H/H	C15,O7,S5
Maintainability	2	3	4	H/M/H	C16,O7,S5
Distribution	4	4	4	H/H/H	C17,O10,S5
Model/provider	3	4	4	M/H/H	C18,O9,S8
Performance	2	3	3	H/M/H	C19,O7,S9
Ecosystem	3	4	4	M/H/H	C20,O10,S10

Bounded uncertainty: exact OMO peer Senpi 2026.8.26-2 was not materialized; remote GitHub exact-head state is not refreshed here; competitor maintainability and remote-operations evidence is asymmetric; model compliance and host implicit skill loading remain unobservable.

Code	Source
S1	devlog/.senpi/packages/coding-agent/src/core/extensions/builtin/goal/store.ts:19-82,148-208; todo-gate.ts:8-30
S2	devlog/.senpi/packages/coding-agent/src/core/dynamic-prompt/intent-gate.ts:14-45; parallel-tools.ts:1-6
S3	devlog/.senpi/packages/agent/src/agent-loop.ts:846-993; devlog/.omo/packages/senpi-task/src/lifecycle/shutdown.ts:9-67
S4	devlog/.senpi/packages/coding-agent/src/core/extensions/builtin/permission-system/index.ts:102-155; core/project-trust.ts:46-95
S5	devlog/.senpi/github/workflows/ci.yml:13-187; build-binaries.yml:81-140
S6	devlog/.senpi/packages/coding-agent/src/core/dynamic-prompt/build.ts:61-101; tool-section.ts:12-45
S7	devlog/.senpi/packages/coding-agent/src/core/extensions/builtin/terminal/monitor-registry.ts:55-164; goal/monitor-continuation.ts:323-618
S8	devlog/.senpi/package.json:15-45; packages/coding-agent/package.json:10-50
S9	devlog/.senpi/github/workflows/perf-trend.yml:1-13,82-88
S10	devlog/.senpi/packages/coding-agent/src/core/extensions/builtin/index.ts:61-100; devlog/.senpi/package.json:5-14